

ILLINOIS DATA SCIENCE INITIATIVE

TECHNICAL REPORTS

Processing Custom Input with Hadoop

Author:
Nishil Shah

February 23, 2017

Processing Custom Input With Hadoop

NISHIL SHAH^{1,3} AND PROFESSOR ROBERT J. BRUNNER^{2,3}

¹National Center For Supercomputing Applications (NCSA)

²Laboratory for Computation, Data, and Machine Learning

³Illinois Data Science Initiative

Compiled March 1, 2017

By implementing our own `FileInputFormat` and `RecordReader` for MapReduce, we can gain control over how Hadoop reads files, constructs input splits, and builds key-value pairs for the mapper.

<https://github.com/lcdm-uiuc>

1. INTRODUCTION

Hadoop includes a collection of input formats (`DBInputFormat`, `KeyValueTextInputFormat`, `NLineInputFormat`, etc...) used to read data files in those respective forms. Typically, Hadoop MapReduce jobs split input data into individual sections, operating on each separately. In some cases, it becomes necessary to customize input splits to avoid reading incomplete data. Also, we may want to read files in obscure formats. We can accomplish this by writing our own `FileInputFormat` and `RecordReader` classes. Specifically, in this report we will read entire files at once.

2. ASSUMPTIONS

This technical report assumes:

- We have the Blockchain downloaded, as shown in "placeholder_report_name".
- The data is stored on HDFS, as explained in "placeholder_report_name".

3. THE DATA

At the time of writing, the Blockchain is stored in a compressed format of over 450,000 binary files. Each file represents one block of the chain. Furthermore, each file contains block metadata as well as a sequence of confirmed transactions. Directly extracting data from a large number of files in this format would be unnecessary difficult and computationally expensive. The files also contain a lot of extraneous information we would like to ignore. Therefore, we will apply MapReduce to transform the Blockchain into a more friendly form.

4. FILEINPUTFORMAT

The `FileInputFormat` class is primarily responsible for creating input splits from files. It also creates a `RecordReader` which builds key-value pairs from each `InputSplit`.

To create a custom `BlockFileInputFormat` we extend `FileInputFormat<K,V>` where `K` and `V` are types of the output key and value pairs. For our purposes, we will use `NullWritable` and

`BytesWritable`, respectively. `NullWritable` reads/writes no bytes; we use it as a placeholder for the key. `BytesWritable`, as the name suggests, stores a sequence of bytes. To prevent splitting the input file, we simply override `isSplittable()` as shown below.

```
@Override
protected boolean isSplittable(JobContext context, Path
    filename) {
    return false;
}
```

We will need to construct our own `RecordReader`. So, we add the following:

```
@Override
public RecordReader<NullWritable, BlockWritable>
    createRecordReader(Input split, TaskAttemptContext
        context) {
    return new BlockFileRecordReader();
}
```

That is all we need to implement in our custom `FileInputFormat` class.

5. RECORDREADER

As explained previously, `RecordReader<K,V>` builds key-value pairs from the input and passes them to the mapper. Let's create our custom `BlockFileRecordReader`. First, we instantiate class variables to hold our input data, generated key-value pair, and a flag:

```
private InputSplit inputSplit;
private Configuration conf;
private boolean processedBlockFile = false;

private NullWritable key = NullWritable.get();
private BytesWritable value = new BytesWritable();
```

In general, the `nextKeyValue()` function is responsible for reading the `InputSplit` (the entire file in our case) to set the class' key and value. As previously mentioned, our key is type `NullWritable` and value is type `BytesWritable`. In `nextKeyValue()`, we read in the bytes of the `InputSplit` and set the value to those

bytes.

We could do further processing on those bytes in this function to extract specific data if we wanted to. Then, we would need to construct custom Writable objects to pass to the mapper. This is explained in detail in "Building Writable Data Types for Hadoop". Instead, we leave this work for the mapper.

Our function looks like this:

```
public boolean nextKeyValue() throws IOException {
    if(!processedBlock) {
        FileSplit fileSplit = (FileSplit)inputSplit;
        Path filePath = fileSplit.getPath();
        FileSystem fileSystem =
            filePath.getFileSystem(conf);

        int splitLength = (int)fileSplit.getLength();
        byte[] blockBytes = new byte[splitLength];

        FSDataInputStream in = null;
        try {
            in = fileSystem.open(filePath);
            IOUtils.readFully(in, blockBytes, 0,
                blockBytes.length);
            value.set(blockBytes, 0, blockBytes.length);
        } finally {
            IOUtils.closeStream(in);
        }
        return processedBlock = true;
    }
    return false;
}
```

To successfully extend RecordReader, we must also override the functions `getCurrentKey()`, `getCurrentValue()`, `getProgress()`, and `close()`. `getProgress()` returns how much of the input the RecordReader has processed (0.0 - 1.0). Since we only use each FileSplit once, we can leave `close()` empty.

6. PUTTING IT ALL TOGETHER

Since our custom FileInputFormat and RecordReader are in separate Java files, we package them by adding `package blockparser;` at the top of each file. To assure the classes are found at runtime, we move them inside a new directory `blockparser` in the project directory.

Finally, to put our custom FileInputFormat to use in our main MapReduce driver, we `import blockparser.BlockFileInputFormat;` and configure the Job as follows:

```
Job job = Job.getInstance(conf, "format blockchain");
//...other configurations
job.setInputFormatClass(BlockFileInputFormat.class);
```

Our mapper will now receive a BytesWritable instance which contains the entire contents of a block file. The map function will then parse the bytes using an external library, *bitcoinj*. From that, we construct BlockWritable and TransactionWritable instances which can be sent to the reducer for final processing. More information on this project can be found in Technical Report *xxxx*.

7. USING EXTERNAL LIBRARIES WITH HADOOP

Luckily, there exists an open source package called *bitcoinj* which we can use to parse the byte contents of each block file. The *bitcoinj* library includes functionality to generate a Block object from this sequence of bytes. From this Block, we can extract just

the information we need into custom data types, BlockWritable and TransactionWritable for use in MapReduce. Refer to Technical Report *xxxx* to learn about creating custom key and value types for Hadoop MapReduce.

To compile and run MapReduce using an external jar, add the jar to the classpath. By running `echo $HADOOP_CLASSPATH` on the command-line, you'll see a list of directories that belong to Hadoop's classpath on your system. Moving the jar (*bitcoinj.jar* in our case) to any of those directories does the trick. In addition, the following configuration should be added to your job in the main driver to ensure the jar is located during runtime.

```
job.addFileToClassPath("path/to/bitcoinj.jar");
```

8. CONCLUSION

Hadoop makes it easy to process data in various formats. By writing custom FileInputFormat and RecordReader classes, we can customize the way Hadoop splits and reads input for use in MapReduce.